


Structured and Deep Similarity Matching via Structured and Deep Hebbian Networks

Dina Obeid, Hugo Ramambason and Cengiz Pehlevan



Harvard John A. Paulson
School of Engineering
and Applied Sciences

Hebbian plasticity implements gradient-based learning

Overview

- 1 We introduce structured and deep similarity matching cost functions.
- 2 We show how they can be optimized in a gradient-based manner by structured and deep neural networks with local learning rules.
- 3 Credit assignment problem is solved elegantly by a factorization of the dual learning objective to synapse specific local objectives.
- 4 Simulations show that our networks learn meaningful features.

Introduction

- How can the brain's synapses, with access to only local information about the network, do gradient-based optimization of an objective function?
- Researchers have been tackling this problem by searching for a biologically-plausible implementation of the backpropagation algorithm, but a fully plausible implementation is not yet available [4].
- We focus on networks already operating with biologically-plausible learning rules. We ask whether one can formulate network-wide learning cost functions for such networks and whether these networks achieve efficient "credit assignment" by performing gradient-based learning.
- We generalize the similarity matching framework [3] to structured and multilayer architectures with feedback connections and generic activation functions.

Regularized similarity matching and gradient-based learning in a Hebbian/anti-Hebbian network

Regularized similarity matching:

Inputs: $\mathbf{x}_1, \dots, \mathbf{x}_T \in \mathbb{R}^K$, and outputs: $\mathbf{r}_1, \dots, \mathbf{r}_T \in \mathbb{R}^N$.

$$\min_{\mathbf{r}_1, \dots, \mathbf{r}_T} \frac{1}{2T^2} \sum_{t=1}^T \sum_{t'=1}^T (\mathbf{x}_t \cdot \mathbf{x}_{t'} - \mathbf{r}_t \cdot \mathbf{r}_{t'})^2 + \frac{2}{T} \sum_{t=1}^T \|\mathbf{F}(\mathbf{r}_t)\|_1 \quad (1)$$

s.t. $a \leq \mathbf{r}_t \leq b$, $t = 1, \dots, T$.

Here, the bounds and the regularization function act elementwise.

Dual problem: Introduce new auxiliary variables $\mathbf{W} \in \mathbb{R}^{N \times K}$ and $\mathbf{L} \in \mathbb{R}^{N \times N}$:

$$\begin{aligned} -\frac{1}{T^2} \sum_t \sum_{t'} \mathbf{x}_t^\top \mathbf{x}_{t'} \mathbf{r}_t^\top \mathbf{r}_{t'} &= \min_{\mathbf{W} \in \mathbb{R}^{N \times K}} -\frac{2}{T} \sum_t \mathbf{x}_t^\top \mathbf{W}^\top \mathbf{r}_t + \text{Tr } \mathbf{W}^\top \mathbf{W} \\ \frac{1}{2T^2} \sum_t \sum_{t'} (\mathbf{r}_t^\top \mathbf{r}_{t'})^2 &= \max_{\mathbf{L} \in \mathbb{R}^{N \times N}} \frac{1}{T} \sum_t \mathbf{r}_t^\top \mathbf{L} \mathbf{r}_t - \frac{1}{2} \text{Tr } \mathbf{L}^\top \mathbf{L}. \end{aligned} \quad (2)$$

A dual min-max formulation of similarity matching:

$$\min_{\mathbf{W} \in \mathbb{R}^{N \times K}} \max_{\mathbf{L} \in \mathbb{R}^{N \times N}} \frac{1}{T} \sum_{t=1}^T l_t(\mathbf{W}, \mathbf{L}, \mathbf{x}_t) \quad (3)$$

where

$$l_t := \text{Tr } \mathbf{W}^\top \mathbf{W} - \frac{1}{2} \text{Tr } \mathbf{L}^\top \mathbf{L} + \min_{\mathbf{r}_t} \left(-2\mathbf{r}_t^\top \mathbf{W} \mathbf{x}_t + \mathbf{r}_t^\top \mathbf{L} \mathbf{r}_t + 2\|\mathbf{F}(\mathbf{r}_t)\|_1 \right). \quad (4)$$

Földiak's Hebbian/anti-Hebbian network [1]:

In the first step, the algorithm minimizes l_t with respect to \mathbf{r}_t by running the neural dynamics (7) until convergence. Minimization is achieved because the argument of min in (4), $E = -2\mathbf{r}_t^\top \mathbf{W} \mathbf{x}_t + \mathbf{r}_t^\top \mathbf{L} \mathbf{r}_t + 2\|\mathbf{F}(\mathbf{r}_t)\|_1$, is a Lyapunov function on the neural dynamics (7), if, within the bounds on the output, the regularizer is related to the neural activation function as:

$$F'(r) = u - r, \quad \text{where } r = f(u). \quad (5)$$

Explicit credit assignment:

$$\sum_{i=1}^N \sum_{j=1}^K \left(-2W_{ij} r_{t,i}^* x_{t,j} + W_{ij}^2 \right) + \sum_{i=1}^N \sum_{j=1}^N \left(L_{ij} r_{t,i}^* r_{t,j}^* - \frac{1}{2} L_{ij}^2 \right). \quad (6)$$

For each input $\mathbf{x} \in \mathbb{R}^K$, run until convergence:

$$\tau \frac{d\mathbf{u}(s)}{ds} = -\mathbf{u}(s) + \mathbf{W} \mathbf{x} - (\mathbf{L} - \mathbf{I}) \mathbf{r}(s), \quad \mathbf{r}(s) = \mathbf{f}(\mathbf{u}(s)). \quad (7)$$

Synaptic updates:

$$\Delta W_{ij} = \eta (r_i x_j - W_{ij}), \quad \Delta L_{ij} = \frac{\eta}{2} (r_i r_j - L_{ij}). \quad (8)$$

Structured Similarity Matching

$$\min_{\substack{a \leq \mathbf{r}_t \leq b, \\ t=1, \dots, T}} \frac{1}{T^2} \sum_{t=1}^T \sum_{t'=1}^T \left(-\sum_{i,j} x_{t,i} x_{t',i} r_{t,j} r_{t',j} c_{ij}^W + \frac{1}{2} \sum_{i,j} r_{t,i} r_{t',i} r_{t,j} r_{t',j} c_{ij}^L \right) + \frac{2}{T} \sum_{t=1}^T \|\mathbf{F}(\mathbf{r}_t)\|_1. \quad (9)$$

Through the choice of c_{ij}^W and c_{ij}^L , one can design many topologies for the interactions between inputs and outputs, and outputs themselves.

Structured and Deep Similarity Matching

Deep Similarity Matching:

For notation convenience, we set $\mathbf{r}_t^{(0)} := \mathbf{x}_t$ and $N^{(0)} := K$, and define deep similarity matching with P layers as:

$$\min_{\substack{a \leq \mathbf{r}_t^{(p)} \leq b, \\ t=1, \dots, T, \\ p=1, \dots, P}} \sum_{p=1}^P \frac{\gamma^{p-P}}{2T^2} \sum_{t=1}^T \sum_{t'=1}^T \left(\mathbf{r}_t^{(p-1)} \cdot \mathbf{r}_{t'}^{(p-1)} - \mathbf{r}_t^{(p)} \cdot \mathbf{r}_{t'}^{(p)} \right)^2 + \sum_{p=1}^P \frac{2\gamma^{p-P}}{T} \sum_{t=1}^T \|\mathbf{F}(\mathbf{r}_t^{(p)})\|_1 \quad (10)$$

where $\gamma \geq 0$ is a parameter and $\mathbf{r}_t^{(p)} \in \mathbb{R}^{N^{(p)}}$.

Structured and deep similarity matching:

$$\begin{aligned} \min_{\substack{a \leq \mathbf{r}_t^{(p)} \leq b, \\ t=1, \dots, T, \\ p=1, \dots, P}} \sum_{p=1}^P \frac{\gamma^{p-P}}{T^2} \sum_{t=1}^T \sum_{t'=1}^T \left(-\sum_{i=1}^{N^{(p-1)}} \sum_{j=1}^{N^{(p)}} r_{t,i}^{(p-1)} r_{t',i}^{(p-1)} r_{t,j}^{(p)} r_{t',j}^{(p)} c_{ij}^{W(p)} \right. \\ \left. + \frac{(1 + \gamma(1 - \delta_{pP}))}{2} \sum_{i=1}^{N^{(p)}} \sum_{j=1}^{N^{(p)}} r_{t,i}^{(p)} r_{t',i}^{(p)} r_{t,j}^{(p)} r_{t',j}^{(p)} c_{ij}^{L(p)} \right) \\ \left. + \sum_{p=1}^P \frac{2\gamma^{p-P}}{T} \sum_{t=1}^T \|\mathbf{F}(\mathbf{r}_t^{(p)})\|_1. \end{aligned} \quad (11)$$

Contact Information

- Email: dinaobeid@seas.harvard.edu
- Web: <http://pehlevan.seas.harvard.edu>
- Code: <https://github.com/Pehlevan-Group>

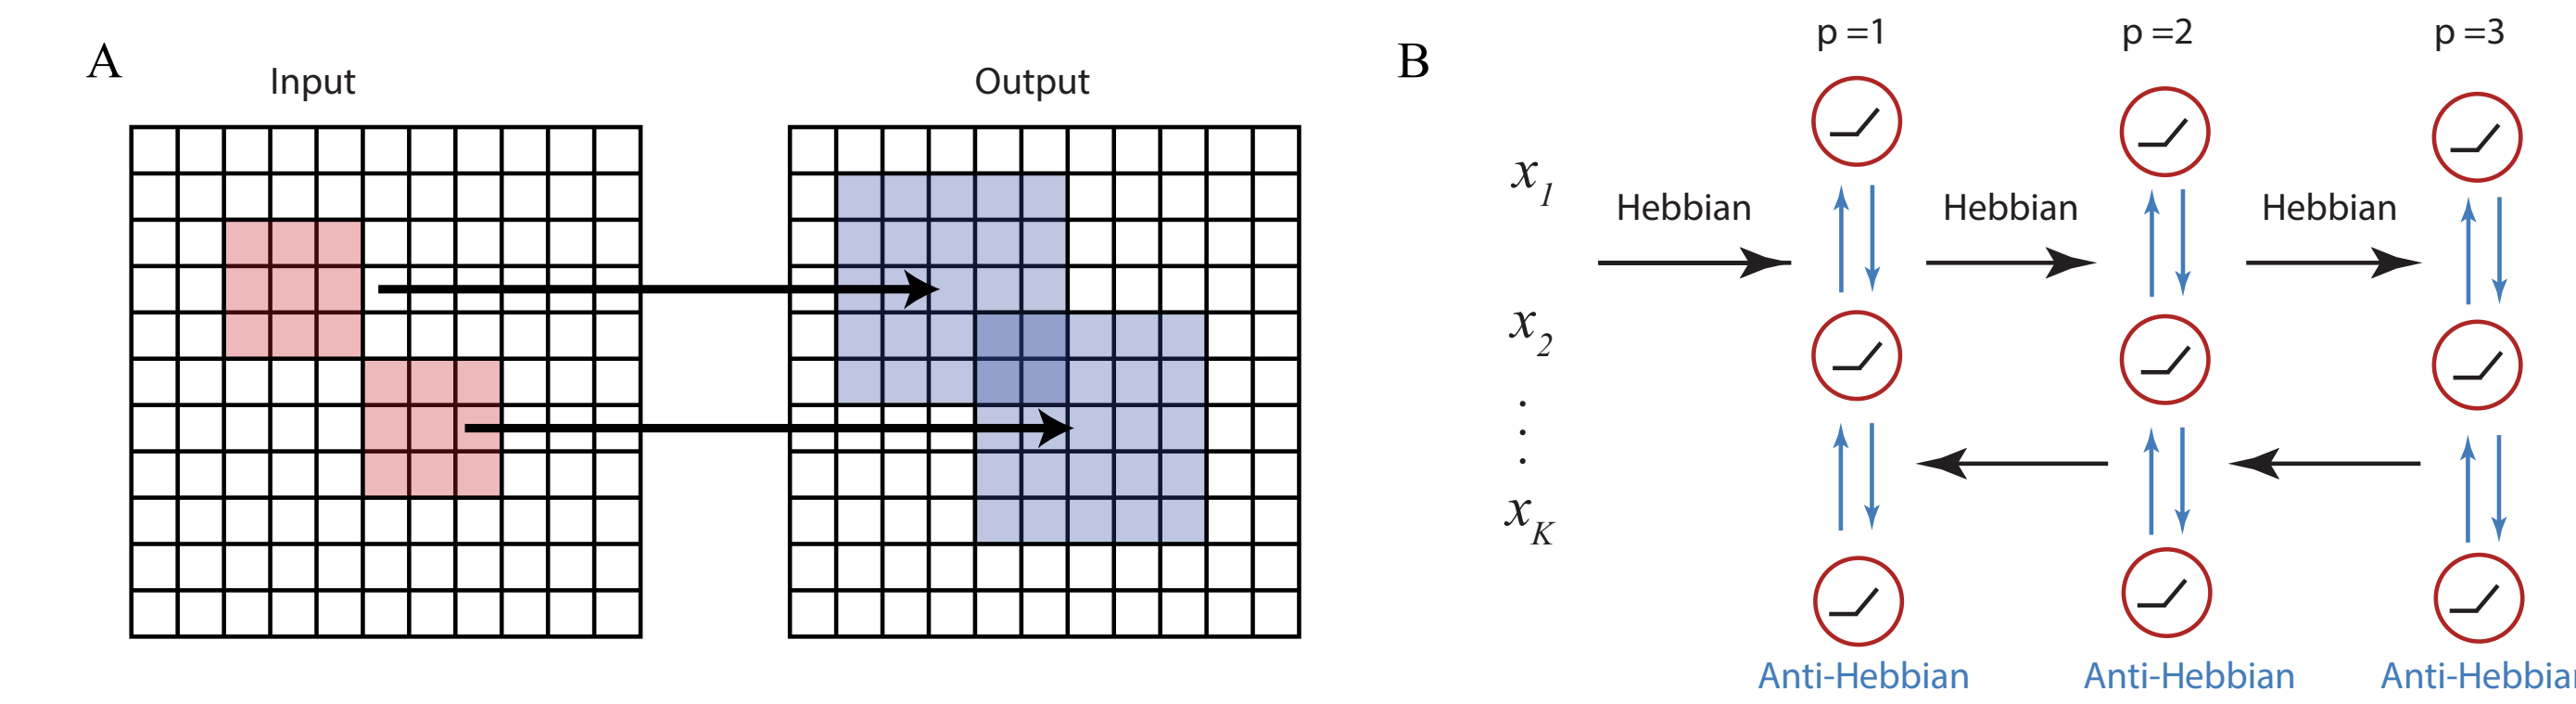


Figure: A) Locally connected similarity matching and structured Hebbian/anti-Hebbian network. B) Deep similarity matching and deep Hebbian/anti-Hebbian network. Arrows illustrate synaptic connections. One can introduce structure for all components of the connectivity.

Structured and Deep Hebbian/anti-Hebbian Networks

We introduce dual variables $W_{ij}^{(p)}$ and $L_{ij}^{(p)}$ for interactions with positive structure constants, define variables

$$\bar{W}_{ij}^{(p)} = \begin{cases} W_{ij}^{(p)}, & c_{ij}^{W(p)} \neq 0 \\ 0, & c_{ij}^{W(p)} = 0 \end{cases} \quad \bar{L}_{ij}^{(p)} = \begin{cases} L_{ij}^{(p)}, & c_{ij}^{L(p)} \neq 0 \\ 0, & c_{ij}^{L(p)} = 0 \end{cases} \quad (12)$$

for notation convenience, and rewrite (11) as

$$\min_{\bar{\mathbf{W}}^{(1)}, \dots, \bar{\mathbf{W}}^{(p)}} \max_{\bar{\mathbf{L}}^{(1)}, \dots, \bar{\mathbf{L}}^{(p)}} \frac{1}{T} \sum_{t=1}^T l_t \left(\bar{\mathbf{W}}^{(1)}, \dots, \bar{\mathbf{W}}^{(p)}, \bar{\mathbf{L}}^{(1)}, \dots, \bar{\mathbf{L}}^{(p)}, \mathbf{r}_t^{(0)} \right) \quad (13)$$

where

$$\begin{aligned} l_t := \sum_{p=1}^P \sum_{\substack{i,j \\ c_{ij}^{W(p)} \neq 0}} \frac{\gamma^{p-P}}{c_{ij}^{W(p)}} W_{ij}^{(p)2} - \sum_{p=1}^P \sum_{\substack{i,j \\ c_{ij}^{L(p)} \neq 0}} \frac{\gamma^{p-P}}{2(1 + \gamma(1 - \delta_{pP})) c_{ij}^{L(p)}} L_{ij}^{(p)2} \\ + \min_{\substack{a \leq \mathbf{r}_t^{(p)} \leq b, \\ p=1, \dots, P}} \sum_{p=1}^P \gamma^{p-P} \left(-2\mathbf{r}_t^{(p)\top} \bar{\mathbf{W}}^{(p)} \mathbf{r}_t^{(p-1)} + \mathbf{r}_t^{(p)\top} \bar{\mathbf{L}}^{(p)} \mathbf{r}_t^{(p)} + 2\|\mathbf{F}(\mathbf{r}_t^{(p)})\|_1 \right) \end{aligned} \quad (14)$$

Neural dynamics

Run the following neural network dynamics until convergence,

$$\begin{aligned} \tau \frac{d\mathbf{u}^{(p)}}{ds} &= -\mathbf{u}^{(p)} + \bar{\mathbf{W}}^{(p)} \mathbf{r}^{(p-1)} - (\bar{\mathbf{L}}^{(p)} - \mathbf{I}) \mathbf{r}^{(p)} + (1 - \delta_{pP}) \gamma \bar{\mathbf{W}}^{(p+1)\top} \mathbf{r}^{(p+1)} \\ \mathbf{r}^{(p)} &= \mathbf{f}(\mathbf{u}^{(p)}), \quad p = 1, \dots, P. \end{aligned} \quad (15)$$

Gradient-based learning and local learning rules

$$\begin{aligned} l_t = \sum_{p=1}^P \sum_{\substack{i,j \\ c_{ij}^{W(p)} \neq 0}} \left(-2W_{ij}^{(p)} r_j^{(p)*} r_i^{(p-1)*} + \frac{\gamma^{p-P}}{c_{ij}^{W(p)}} W_{ij}^{(p)2} \right) \\ - \sum_{p=1}^P \sum_{\substack{i,j \\ c_{ij}^{L(p)} \neq 0}} \left(-L_{ij}^{(p)} r_j^{(p)*} r_i^{(p)*} + \frac{\gamma^{p-P}}{2(1 + \gamma(1 - \delta_{pP})) c_{ij}^{L(p)}} L_{ij}^{(p)2} \right). \end{aligned} \quad (16)$$

Local learning rules are derived from the above equation by taking derivatives:

$$\begin{aligned} \Delta W_{ij}^{(p)} &= \eta \gamma^{p-P} \left(r_j^{(p)} r_i^{(p-1)} - \frac{W_{ij}^{(p)}}{c_{ij}^{W(p)}} \right), \\ \Delta L_{ij}^{(p)} &= \frac{\eta}{2} \gamma^{p-P} \left(r_j^{(p)} r_i^{(p)} - \frac{L_{ij}^{(p)}}{(1 + \gamma(1 - \delta_{pP})) c_{ij}^{L(p)}} \right). \end{aligned} \quad (17)$$

Simulations

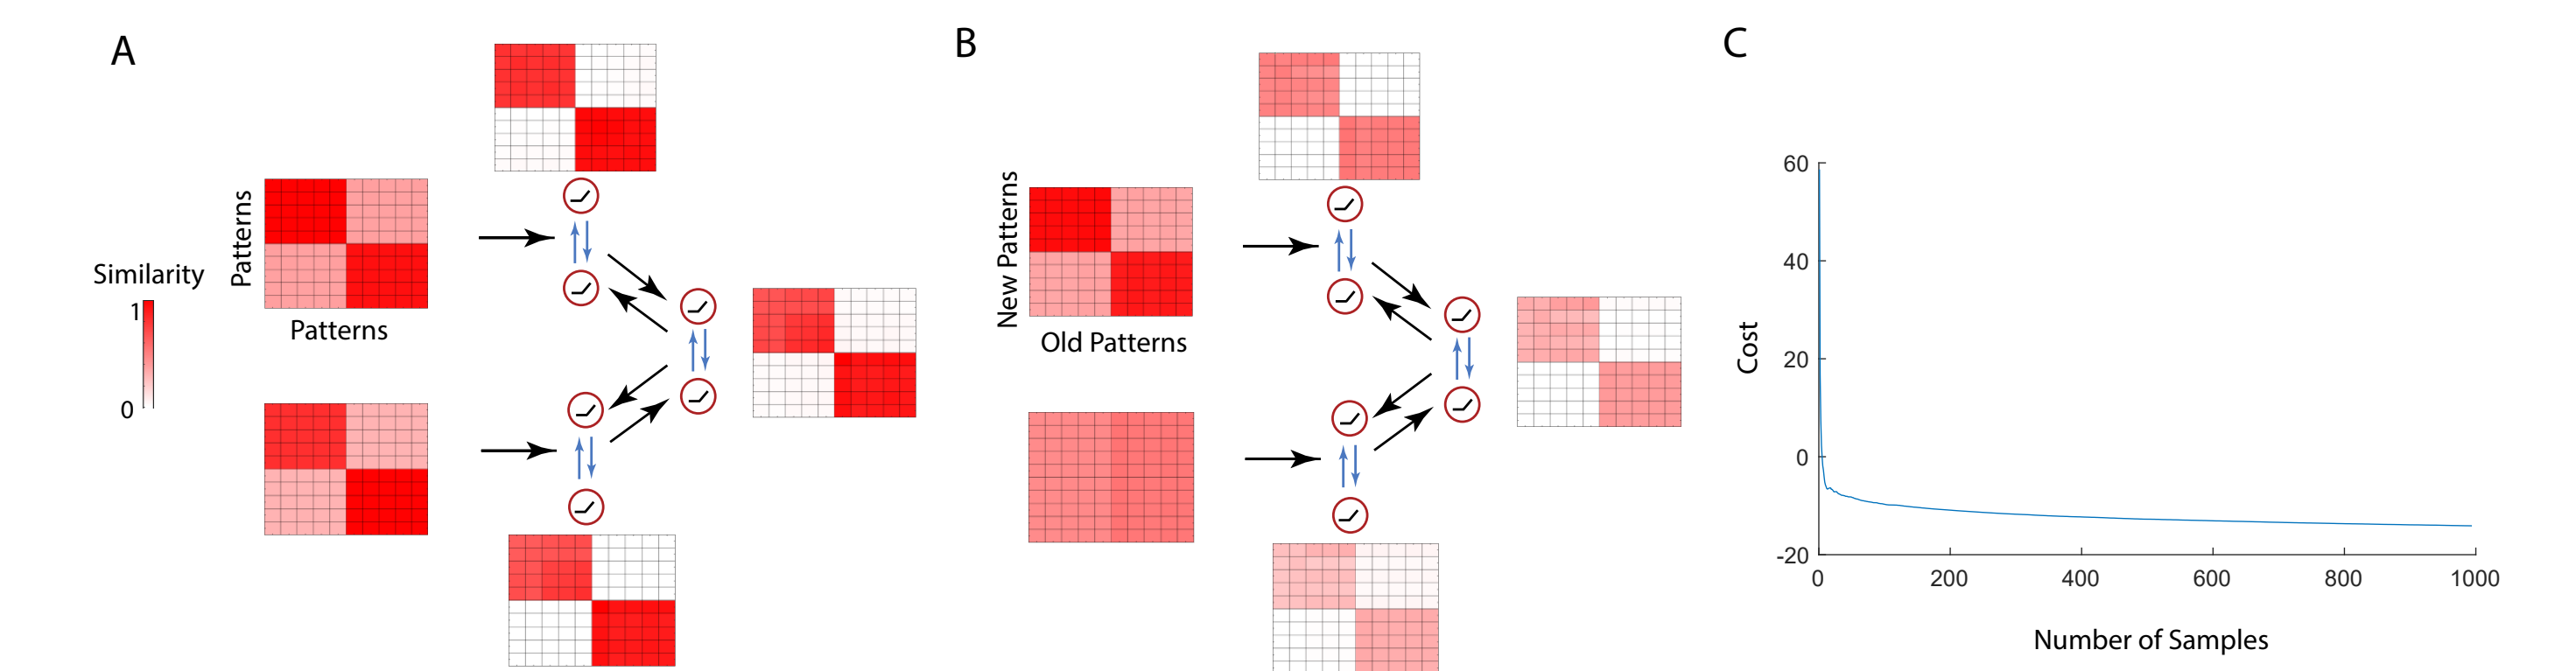


Figure: A two-layer Hebbian/anti-Hebbian network with feedback. For each subnetwork, representational similarity matrices are shown. Similarities are calculated by taking pairwise dot products of patterns and normalizing the largest dot product to 1. A) Network simulated with patterns from a set generated from the same distribution as the training set. B) Network simulated with patterns to the bottom first layer generated from a different distribution. C) Structured and deep similarity matching cost decreases over training.




Figure: Features learned by a 3-layer, locally connected Hebbian/anti-Hebbian neural network on the labeled faces in the wild dataset [2]. Features are calculated by reverse correlation on the dataset, and masking these features to keep only the portions of the dataset which elicits a response in the neuron.

NPS	4	8	16	32	64	100
Classification error (%)	3.87	2.41	1.73	1.60	1.47	1.40

Table: Classification on MNIST data set: we show how the test error decreases as the number of neurons per site (NPS) increases.

References

- [1] P. Földiak. Forming sparse representations by local anti-hebbian learning. *Biological cybernetics*, 64(2):165–170, 1990.
- [2] E. Learned-Miller, G. B. Huang, A. RoyChowdhury, H. Li, and G. Hua. Labeled faces in the wild: A survey. In *Advances in face detection and facial image analysis*, pages 189–248. Springer, 2016.
- [3] C. Pehlevan and D. B. Chklovskii. Neuroscience-inspired online unsupervised learning algorithms. *IEEE Signal Processing Magazine*, 36:88–96, 2019.
- [4] J. C. Whittington and R. Bogacz. Theories of error back-propagation in the brain. *Trends in cognitive sciences*, 2019.

Acknowledgements

This work was funded by Intel Corporation.